

Do agent routers need to look inside the model?

A preliminary audit of internal-state failure probes

Shlok Channawar

September 2026

1 Motivation

Hera, MTRouter and CEB all need an estimate of the acting model’s state *at each step* in order to route or intervene, and all three reconstruct it from outside the model. Hera observes that device/cloud step agreement correlates with predictive entropy—an internal signal—then trains a separate 500M network to read state text instead. MTRouter cannot inspect a mostly-closed model pool, so it learns a per-model residual vector as a proxy for behavioural character it cannot observe. CEB keeps the model frozen and routes through what it can articulate; its Appendix F reports TF-IDF beating dense embeddings, because the retrieval target is near-duplicate actions on near-identical pages. Three workarounds for one missing measurement.

Does that signal exist inside the acting model at each step, is it richer than entropy, and does it generalise along failure *structure* rather than surface similarity? Ruan et al. (2607.06503) is the one paper here that reads hidden states directly: AUC 0.81 for predicting eventual failure from a layer-28 state at the final token of the agent’s first action (TextCraft, Qwen3-1.7B), used to gate an abort cascade. If it means what it appears to, it is the measurement the other three approximate.

2 Preliminary experiment: does the reported metric measure step-level state?

Any per-step decision requires knowing that *this rollout* is going wrong, not that *this task* is hard. Ruan’s reported statistic is a pooled cross-fitted AUC over all episodes, which cannot separate these: a probe that only reads task difficulty scores well on it while carrying no per-rollout information at all.

We replicated the TextCraft/Qwen3-1.7B cell (1000 episodes; 125 depth-one tasks $\times$ 8 rollouts; 20-round cap; L2 logistic probe, $C=1$; task-grouped stratified cross-fitting) and obtained **pooled AUC 0.791** against their 0.81. We then added the control the paper does not report: **within-task AUC**, which compares failures against successes *on the same task*, so task difficulty is held constant by construction.

	pooled AUC	within-task AUC
Hidden state, layer 28 (reported config)	0.791	0.547
Five cheap surface features	0.698	0.615
TF-IDF text + surface	0.717	0.606

99.6% of the comparisons behind the pooled number are between different tasks (173,748 of 174,375); only 627 hold the task fixed. 77% of score variance is between tasks: the probe emits approximately one score per task (Fig. 1). Held to the same task, it falls to chance, and the cheap surface baseline beats it.




Figure 1: **Left:** for each mixed-outcome task, mean probe score on its failures against its successes. A probe carrying within-task information would sit above the dashed diagonal; only 24 of 57 tasks do. **Right:** within-task AUC by layer. The reported layer sits on the downslope; layer 12 wins at both rounds, and at round 1 it wins on the paper’s own selection criterion (pooled AUC 0.830 vs 0.781).

3 What the probe was actually detecting

Classifying all 225 failures by behaviour makes this concrete. At the reported configuration the probe catches **77%** of failures on tasks whose recipe names a count-one plural ingredient (`craft 4 acacia planks using 1 acacia logs`: the plural form is a distinct item and the model silently “corrects” it) and **29%** of every other failure. It is largely a detector for one lexical quirk—the same surface-similarity behaviour that motivates the concern about CEB’s TF-IDF retrieval.

4 The signal exists, but not at the reported coordinates

Sweeping rounds 1–6 and layers $\{4, 8, \dots, 28\}$ (each layer captured in one forward pass) locates it elsewhere. At **layer 12, round 3**, within-task AUC reaches **0.801** (95% CI $[0.74, 0.86]$), against 0.749 at layer 28 and **0.725** for a text model given every prompt, response and error message through round 3. Detection also generalises: looping rises from $\sim 29\%$ to 73% across 213 episodes, format errors to 80%, explicit give-ups to 83% (Fig. 2), while eventual successes stay at 0.195.




Figure 2: Detection by failure mode: from one lexical quirk to failure structure. The reported configuration finds the lexical trap; the better-placed probe finds failure shapes that share no vocabulary.

5 Opportunity, and what would falsify this

The three sub-questions have provisional answers: the signal *does* exist per step, is modestly richer than entropy-like surface features, and generalises along failure structure rather than surface similarity—but only at a readout point the literature does not report, and the reported configuration is precisely the one at which the answer looks negative.

Scope. This audits Ruan’s evaluation methodology. We have not tested whether Hera’s agreement signal or MTRouter’s outcome estimator conflate difficulty with instance; they are motivating contrast, not targets of critique.

Limits. Layer 12 / round 3 is a cell selected from a 7×6 search, so 0.801 is an exploratory maximum, not an effect size. Sweeps ran on MPS against CUDA canonical features (0.547 vs 0.585 at the same cell). The text baseline is TF-IDF, a weak stand-in for an LLM judge. One environment, one 1.7B model, and “failure” means the item was not crafted, not that an unsafe action was taken.

Next, ranked. (i) An LLM-as-judge baseline—the comparison that decides whether internals are worth instrumenting at all. (ii) A pre-registered replication with layer 12 / round 3 fixed in advance on a fresh seed. (iii) A paired fold-level test rather than overlapping intervals. (iv) The same decomposition where the label is an unsafe action rather than a failed task, which is what would make this safety work rather than reliability work.